

Principles for AI Adoption and Governance

1. AI Adoption Is an Organizational Change Process

AI adoption should not be treated only as a software deployment or tool-selection exercise. Sustainable adoption depends on understanding how work is performed, where friction occurs, what knowledge people rely on, and how changes affect roles, workload, trust, incentives, and decision-making.

The central question is not simply, “Where can AI be used?” It is:

Where can AI safely support human work, under what conditions, with what evidence, and with what organizational safeguards?

AI adoption is therefore a combination of:

- workflow analysis;
- organizational learning;
- knowledge management;
- change management;
- risk management;
- human oversight.

2. Begin With Real Workflow Evidence

Adoption planning should start with evidence from actual work rather than generic assumptions about AI capabilities.

Useful evidence may include:

- interviews and interaction notes;
- observed workflow bottlenecks;
- recurring manual tasks;
- staff feedback;
- service or quality issues;
- existing templates, policies, and procedures;

- examples of exceptions and failures;
- workload and coordination problems.

Potential AI opportunities should remain hypotheses until they have been validated with the people who perform or supervise the work.

Evidence should be traceable to its source so reviewers can understand why an opportunity, risk, or recommendation was recorded.

3. Decompose Workflows Before Selecting Technology

A workflow should be divided into smaller steps before deciding which parts can be supported by technology.

Different steps may require different approaches, including:

- generative AI;
- search or retrieval;
- structured templates;
- browser or process automation;
- OCR or document processing;
- deterministic business rules;
- human judgment.

Not every inefficient step requires generative AI. Tool selection should follow workflow analysis rather than precede it.

A useful decision process is:

Understand the workflow

→ divide it into steps

→ identify the decision or information required at each step

→ determine the appropriate tool

→ define human-review boundaries

→ test the revised workflow

4. Start With Small, Repeatable Improvements

Initial adoption efforts should focus on narrowly scoped, repeatable workflow improvements rather than large autonomous systems.

Suitable early opportunities often have:

- a clearly defined task;
- frequent repetition;
- visible inefficiency;
- available source knowledge;
- reversible outcomes;
- limited privacy or safety impact;
- an identifiable reviewer;
- measurable baseline performance.

A small pilot can reveal missing knowledge, quality problems, staff concerns, and implementation constraints before a broader rollout.

Pilot plans should include:

- scope;
- responsible roles;
- required knowledge;
- review checkpoints;
- success measures;
- known risks;
- stop, revise, or escalation criteria.

5. Treat Human Review as Part of the Workflow

Human oversight should be explicitly designed rather than added as a general disclaimer.

Review is especially important for work that is:

- externally visible;

- private or confidential;
- quality-sensitive;
- judgment-heavy;
- consequential for individuals;
- difficult to reverse;
- dependent on incomplete or ambiguous context.

For each AI-supported task, define:

- who reviews the output;
- what the reviewer must check;
- what source information should be consulted;
- which errors require escalation;
- when the output must not be used;
- who remains accountable for the final decision.

AI-generated content should be treated as a draft, recommendation, or analytical input unless it has passed the required review process.

6. Document Failure Cases and Exceptions

Governance should include systematic documentation of what can go wrong.

A failure record should capture:

- what failed or may fail;
- why it failed;
- what context was missing;
- the impact of the failure;
- whether human review could detect it;
- the prevention or mitigation step;
- the related workflow;

- the evidence source.

Failure cases are not merely defects. They help define:

- human-review requirements;
- system boundaries;
- knowledge gaps;
- escalation rules;
- unsuitable use cases;
- future training needs.

Repeated failures should strengthen the case for revising or stopping a workflow rather than being treated as isolated incidents.

7. Build Reliable Knowledge Foundations

AI-supported workflows depend on structured, searchable, and reliable knowledge sources.

Relevant sources may include:

- policies;
- procedures;
- templates;
- frequently asked questions;
- prior decisions;
- training materials;
- past correspondence;
- meeting notes;
- reports;
- public information;
- experienced staff knowledge.

For each proposed workflow, determine:

- what information the system requires;
- where that information is stored;
- who owns it;
- whether it is current;
- whether it is approved for use;
- whether access is permitted;
- what happens when sources disagree;
- how updates are maintained.

A lack of reliable knowledge should be treated as an adoption constraint, not as a problem that prompting alone can solve.

8. Preserve Source Context and Traceability

AI-derived findings should retain enough context for later review.

Useful provenance fields include:

- source record or interaction;
- evidence excerpt;
- reviewer;
- date identified;
- last confirmed date;
- number of supporting observations;
- approval or validation status;
- confidence level;
- related workflow or risk.

Traceability helps distinguish evidence-based findings from unsupported model suggestions.

It also supports:

- audits;

- knowledge review;
- dispute resolution;
- stale-information checks;
- comparison across cases;
- future evaluation.

9. Manage Accumulated Knowledge as Reviewable Evidence

AI-generated workflow knowledge should not automatically be treated as validated truth.

Accumulated findings require periodic human review to determine whether they are:

- reliable;
- outdated;
- duplicated;
- contradicted by newer evidence;
- specific to one context;
- reusable across contexts;
- suitable for formal policy or guidance.

Knowledge may be organized into several layers:

Context-specific knowledge

Evidence about one team, department, workflow, or organization.

Reusable method knowledge

Reviewed principles, procedures, and playbooks that apply across contexts.

Searchable evidence knowledge

Source-grounded records used for retrieval, comparison, or analysis.

Planning knowledge

Reviewed inputs used to create roadmaps, pilot plans, training plans, and success measures.

Clear promotion criteria should govern when an observation becomes an approved reusable practice.

10. Avoid Duplicate Findings Without Losing Evidence

Similar findings from multiple observations should usually be consolidated rather than presented as separate conclusions.

A consolidated record should preserve:

- a normalized title or key;
- all supporting sources;
- evidence excerpts;
- evidence count;
- first-seen and last-seen dates;
- current validation status.

Multiple independent observations should increase the strength of evidence. Deduplication should not erase disagreement, important exceptions, or contextual differences.

11. Include Human-System Risks in Governance

AI adoption can fail even when the technology works.

Governance should examine risks such as:

- workload increases;
- low staff trust;
- unclear role boundaries;
- training gaps;
- knowledge-sharing gaps;
- privacy or quality concerns;
- weak leadership support;

- unclear evaluation criteria;
- insufficient recognition for improvement work;
- incentives that discourage responsible experimentation;
- efficiency gains being converted only into additional workload.

These risks should be assessed alongside technical risks.

A human-system risk record may include:

- risk type;
- description;
- severity;
- affected roles;
- evidence source;
- proposed mitigation;
- responsible owner;
- review date.

12. Align Workload, Evaluation, and Recognition

Efficiency improvements can create resistance when employees believe that using AI will result only in greater workload, reduced job security, or loss of recognition.

Organizations should explicitly discuss how AI-enabled changes affect:

- workload expectations;
- performance evaluation;
- job responsibilities;
- role clarity;
- quality standards;
- professional judgment;
- credit for process improvements;

- opportunities for development.

Employees who responsibly improve workflows should receive appropriate recognition. Adoption programs should not reward only output volume while ignoring quality, learning, risk reduction, or knowledge sharing.

13. Develop AI Capability Through Workflow Learning

AI training should focus on changing real workflows, not only on teaching tools or prompt techniques.

Staff should learn to:

- identify inefficient work;
- break workflows into steps;
- distinguish AI-suitable tasks from unsuitable ones;
- locate required knowledge;
- design review checkpoints;
- test small improvements;
- document failures;
- measure outcomes;
- share reusable examples.

A practical learning cycle is:

```
Identify a small workflow problem
→ test an AI-assisted improvement
→ review the result
→ document what worked and failed
→ share the example
→ store reusable knowledge
→ identify the next improvement
```

Experienced employees should be involved because they are often best positioned to judge quality, detect missing context, and identify exceptions.

14. Establish Clear Safety Boundaries

Each AI initiative should define what the system is and is not permitted to do.

Possible boundaries include:

- no autonomous external communication;
- no automatic implementation of recommendations;
- no access to unapproved systems or repositories;
- no use of confidential information without authorization;
- no automated decisions in sensitive areas;
- no unsupported claims about people, services, or organizational conditions;
- no action without a designated human reviewer.

Technical controls should reinforce these boundaries. Labels such as “draft” or “recommendation” are useful, but they are not substitutes for access controls, approval workflows, logging, and accountability.

15. Protect Privacy and Control Data Access

Governance should identify:

- what data is being processed;
- whether the data is personal, confidential, or regulated;
- where the data is stored;
- which systems receive it;
- who can access it;
- how long it is retained;
- whether it is used for training;
- how it is deleted;
- whether third-party services are involved.

The principle of least privilege should apply. AI systems should receive only the information required for the specific task.

Local or isolated processing may reduce some risks, but it does not eliminate the need for access control, retention rules, secure storage, and review.

16. Separate Advice From Action

AI systems should clearly distinguish between:

- generating a draft;
- recommending an action;
- approving an action;
- executing an action.

Increasing autonomy should require stronger governance.

A useful progression is:

Information retrieval

- draft generation
- recommendation
- human-approved execution
- limited automation with monitoring

Organizations should not move to a higher level of autonomy merely because a lower-risk demonstration was successful.

17. Define Evaluation Before Scaling

A pilot should have baseline measures, target outcomes, and a defined review method.

Possible measures include:

- time required;
- error rate;
- revision rate;
- reviewer effort;
- task completion rate;
- user satisfaction;
- service quality;
- number and severity of exceptions;

- privacy or policy incidents;
- knowledge gaps identified;
- staff adoption and trust;
- sustainability over time.

Evaluation should consider both productivity and quality. A faster workflow is not an improvement if it increases errors, hides risk, transfers work to reviewers, or degrades employee or customer experience.

18. Recognize Evaluation Uncertainty

Many AI adoption decisions do not have a stable or objective ground truth.

The quality of an AI recommendation may depend on:

- the available evidence;
- the completeness of source knowledge;
- the model or system used;
- reviewer expertise;
- organizational context;
- changing AI capabilities.

Outputs should therefore be represented as reviewable working conclusions rather than fixed best practices.

Future evaluation methods may include:

- expert review;
- structured reviewer scoring;
- comparison with human-created outputs;
- regression tests;
- incident analysis;
- user feedback;
- periodic revalidation.

19. Define Governance Roles and Accountability

An AI governance process should identify responsible roles for:

- workflow ownership;
- data ownership;
- source approval;
- technical operation;
- risk review;
- privacy and security;
- output review;
- pilot approval;
- incident response;
- final business decisions.

Accountability should remain with identifiable people or organizational bodies. It should not be transferred to the AI system or described vaguely as “human in the loop.”

20. Conduct Periodic Adoption Reviews

AI-supported workflows should be reviewed after implementation because workflows, personnel, policies, data, and model capabilities change.

Periodic reviews should consider:

- whether the original need still exists;
- whether the workflow remains accurate;
- whether source knowledge is current;
- whether review effort has increased;
- whether new failure cases have appeared;
- whether staff concerns have changed;
- whether usage has expanded beyond the approved scope;

- whether performance remains acceptable;
- whether the workflow should continue, change, or stop.

Review dates and ownership should be established before deployment.

Governance Checklist

Before approving an AI-supported workflow, confirm that:

1. The workflow problem is supported by real evidence.
2. The workflow has been decomposed into clear steps.
3. AI is appropriate for the selected steps.
4. Required knowledge sources are identified and reliable.
5. Data access and privacy requirements are documented.
6. Known failure cases and exceptions are recorded.
7. Human-review responsibilities are explicit.
8. Accountability for final decisions is assigned.
9. Staff workload, role, trust, and incentive impacts have been considered.
10. Pilot scope and stop criteria are defined.
11. Baseline and success measures are established.
12. Outputs remain traceable to source evidence.
13. Technical controls enforce the stated safety boundaries.
14. Accumulated findings will be periodically reviewed.
15. Scaling requires additional evidence and governance approval.

Core Operating Principle

AI adoption should progress through evidence, limited experimentation, human review, organizational learning, and measured expansion.

Observe real work

- identify a bounded opportunity
- verify knowledge and risks
- design human oversight
- run a limited pilot
- evaluate outcomes and impacts
- document learning
- revise, stop, or scale